

Gaussian Processes

Advanced Statistical Inference

Simone Rossi

Multivariate Gaussian: Marginalization and Conditioning

Let $\mathbf{z} = (z_1, z_2, z_3)^\top \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ with

$$\boldsymbol{\mu} = \begin{pmatrix} 1 \\ 2 \\ 0 \end{pmatrix}, \quad \boldsymbol{\Sigma} = \begin{pmatrix} 4 & 2 & 0 \\ 2 & 3 & 1 \\ 0 & 1 & 2 \end{pmatrix}.$$

Partition the vector as $\mathbf{z} = (\mathbf{z}_a^\top, z_b)^\top$ where $\mathbf{z}_a = (z_1, z_2)^\top$ and $z_b = z_3$, so that

$$\boldsymbol{\mu}_a = \begin{pmatrix} 1 \\ 2 \end{pmatrix}, \quad \mu_b = 0, \quad \boldsymbol{\Sigma}_{aa} = \begin{pmatrix} 4 & 2 \\ 2 & 3 \end{pmatrix}, \quad \boldsymbol{\Sigma}_{ab} = \begin{pmatrix} 0 \\ 1 \end{pmatrix}, \quad \Sigma_{bb} = 2.$$

1. **Marginalization.** State the general rule for the marginal $p(\mathbf{z}_a)$ of a partitioned Gaussian. Apply it to compute $p(\mathbf{z}_a) = p(z_1, z_2)$ explicitly, giving its mean vector and covariance matrix. Analyze all the dimensions of the matrices and vectors involved to verify that they are consistent.
2. **Conditioning.** State the general formula for the conditional $p(\mathbf{z}_a | z_b)$:

$$p(\mathbf{z}_a | z_b) = \mathcal{N} \left(\underbrace{\boldsymbol{\mu}_a + \boldsymbol{\Sigma}_{ab} \Sigma_{bb}^{-1} (z_b - \mu_b)}_{\boldsymbol{\mu}_{a|b}}, \underbrace{\boldsymbol{\Sigma}_{aa} - \boldsymbol{\Sigma}_{ab} \Sigma_{bb}^{-1} \boldsymbol{\Sigma}_{ab}^\top}_{\boldsymbol{\Sigma}_{a|b}} \right).$$

For the observation $z_b = 2$, compute numerically:

- the conditional mean $\boldsymbol{\mu}_{a|b}$;
 - the conditional covariance $\boldsymbol{\Sigma}_{a|b}$.
3. **Interpretation.** Compare $\boldsymbol{\Sigma}_{aa}$ and $\boldsymbol{\Sigma}_{a|b}$. Which is “larger”? What does this tell you about the effect of observing z_b on our uncertainty about $\mathbf{z}_a$? Which component (z_1 or z_2) is affected more, and why?

From Parametric to Non-Parametric Models

1. In Bayesian linear regression we place a prior over weights $\mathbf{w}$ and obtain a posterior $p(\mathbf{w} \mid \mathbf{y}, \mathbf{X})$. Explain in words what it means to instead place a prior directly over functions $f : \mathcal{X} \rightarrow \mathbb{R}$.
2. A Gaussian process is defined as: a collection of random variables such that every finite subset has a joint Gaussian distribution. Let $f \sim \mathcal{GP}(m, k)$ where $m(\mathbf{x}) = 0$ and $k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2}\|\mathbf{x} - \mathbf{x}'\|^2\right)$.
 - Write the joint distribution $p(f(\mathbf{x}_1), f(\mathbf{x}_2))$ for two inputs $\mathbf{x}_1, \mathbf{x}_2$.
 - What is the marginal distribution of $f(\mathbf{x}_1)$?

Kernel Functions

1. The squared exponential (RBF) kernel is

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right).$$

- What does the lengthscale ℓ control? What happens as $\ell \rightarrow 0$ and $\ell \rightarrow \infty$?
 - What does the signal variance σ_f^2 control?
2. For inputs $x_1 = 0, x_2 = 1, x_3 = 3$ with $\sigma_f^2 = 1$ and $\ell = 1$, compute the 3×3 kernel matrix $\mathbf{K}$ where $K_{ij} = k(x_i, x_j)$.
 3. A valid kernel must be symmetric and positive semi-definite. Verify that the RBF kernel is symmetric. Why is positive semi-definiteness important for the covariance matrix of a GP?
 4. Consider the linear kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}'$. What kind of functions does a GP with this kernel represent?

GP Regression (Exact Inference)

Consider a noise-corrupted observation model:

$$y_n = f(\mathbf{x}_n) + \varepsilon_n, \quad \varepsilon_n \sim \mathcal{N}(0, \sigma_n^2),$$

with prior $f \sim \mathcal{GP}(0, k)$.

1. Write the joint distribution of the observations $\mathbf{y} = (y_1, \dots, y_N)^\top$ and the latent function values $\mathbf{f} = (f(\mathbf{x}_1), \dots, f(\mathbf{x}_N))^\top$.

2. The joint distribution of training outputs $\mathbf{y}$ and test function values $\mathbf{f}_\star = f(\mathbf{X}_\star)$ is Gaussian:

$$\begin{pmatrix} \mathbf{y} \\ \mathbf{f}_\star \end{pmatrix} \sim \mathcal{N}\left(\mathbf{0}, \begin{pmatrix} \mathbf{K} + \sigma_n^2 \mathbf{I} & \mathbf{K}_\star^\top \\ \mathbf{K}_\star & \mathbf{K}_{\star\star} \end{pmatrix}\right),$$

where $\mathbf{K} = k(\mathbf{X}, \mathbf{X})$, $\mathbf{K}_\star = k(\mathbf{X}_\star, \mathbf{X})$, $\mathbf{K}_{\star\star} = k(\mathbf{X}_\star, \mathbf{X}_\star)$. Derive the posterior predictive mean and covariance:

$$\mu_\star = \mathbf{K}_\star (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}, \quad \Sigma_\star = \mathbf{K}_{\star\star} - \mathbf{K}_\star (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_\star^\top.$$

3. Suppose we have a single training point $x_1 = 0$ with $y_1 = 1$, and we want to predict at $x_\star = 1$. Use the RBF kernel with $\sigma_f^2 = 1$, $\ell = 1$, and $\sigma_n^2 = 0.1$.
- Compute $K = k(x_1, x_1)$, $k_\star = k(x_\star, x_1)$, $k_{\star\star} = k(x_\star, x_\star)$.
 - Compute the posterior predictive mean $\mu_\star$ and variance $\sigma_\star^2$.
4. **Two training points.** Now suppose we have two training points $x_1 = 0$, $y_1 = 1$ and $x_2 = 2$, $y_2 = -1$, and we want to predict at $x_\star = 1$. Use the RBF kernel with $\sigma_f^2 = 1$, $\ell = 1$, $\sigma_n^2 = 0.25$.
- Build the 2×2 matrix $\mathbf{K} + \sigma_n^2 \mathbf{I}$ and compute its inverse.
 - Compute the posterior predictive mean $\mu_\star$ and variance $\sigma_\star^2$ at $x_\star = 1$.
5. **Effect of noise.** Repeat the single-training-point prediction from the previous section ($x_1 = 0$, $y_1 = 1$, $x_\star = 1$, RBF with $\sigma_f^2 = 1$, $\ell = 1$) for two noise levels: $\sigma_n^2 = 0.01$ and $\sigma_n^2 = 4$.
- Compute $\mu_\star$ and $\sigma_\star^2$ for each case.
 - Describe qualitatively how increasing σ_n^2 affects the posterior mean and variance at the test point, and give an intuitive explanation.

Hyperparameter Optimisation

1. The marginal likelihood (model evidence) for GP regression is

$$\log p(\mathbf{y} \mid \mathbf{X}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{y}^\top (\mathbf{K}_\theta + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K}_\theta + \sigma_n^2 \mathbf{I}| - \frac{N}{2} \log 2\pi.$$

Identify the three terms and give an intuitive explanation of the role each plays in selecting the kernel hyperparameters $\boldsymbol{\theta} = (\ell, \sigma_f^2, \sigma_n^2)$.

2. Why is maximising the marginal likelihood preferable to maximising the training data likelihood $p(\mathbf{y} \mid \mathbf{X}, \mathbf{f})$ for choosing hyperparameters?
3. What is the computational complexity of exact GP regression in terms of the number of training points N ? What bottleneck causes this? Name one strategy to reduce this cost.